ABSTRACT

We consider the classical geometric problem of determining a shortest path through a weighted domain. We present approximation algorithms that compute ε-*short* paths, i.e., paths whose costs are within a factor of $1 + \varepsilon$ of the shortest path costs, for an arbitrary constant $\varepsilon > 0$, for the following geometric configurations:

SPPS Problem: We are given a polyhedron $\mathcal{P}$ consisting of n convex faces[1] and each face has a positive non-zero real valued weight. The *shortest path on polyhedral surface problem* (SPPS) is to compute a path of least cost that remains on the surface of $\mathcal{P}$ between any two vertices, where the cost of the path is defined to be the weighted sum of Euclidean lengths of the sub-paths within each face. Our algorithm runs in $O(\frac{n}{\varepsilon} \log \frac{1}{\varepsilon}(\frac{1}{\sqrt{\varepsilon}} + \log n))$ time for $0 < \varepsilon < 1$. The run time improves to $O(n \log n)$ for $\varepsilon \geq 1$ and to $O(\frac{n}{\varepsilon} \log \frac{1}{\varepsilon} \log n)$ when all weights are equal.

WRP-3D Problem: We are given a subdivision of $\Re^3$ consisting of n convex regions. Each face has associated with it a positive non-zero real valued weight. The *shortest path problem in three dimensions* (SP3D) is to compute a path of least cost between any two vertices, where the cost of the path is defined to be the weighted sum of Euclidean lengths of the sub-paths within each region. We present the first polynomial time approximation scheme. Our algorithm runs in $O(\frac{n}{\varepsilon^3} \log \frac{1}{\varepsilon}(\frac{1}{\sqrt{\varepsilon}} + \log n))$ time. The run time improves to $O(\frac{n}{\varepsilon^3} \log \frac{1}{\varepsilon} \log n))$ when all weights are equal. This can be used to solve the shortest path problem amidst obstacles in 3-dimensional Euclidean space (ESP-3D).

---

[1] For ease of notation we assume throughout the SPPS portion of this paper that the number of faces is proportional to the number of vertices although our algorithms are stated in the general sense.

Approximation algorithms for geometric shortest path problems[*]

Lyudmil Aleksandrov
Bulgarian Academy of
Sciences
1113 Sofia, Bulgaria.
lyualeks@argo.bas.bg

Anil Maheshwari
School of Computer Science
Carleton University
Ottawa, Canada K1S 5B6.
anil@scs.carleton.ca

Jörg-Rüdiger Sack
School of Computer Science
Carleton University
Ottawa, Canada K1S 5B6.
sack@scs.carleton.ca

1. INTRODUCTION

1.1 Motivation

Shortest path problems are among the fundamental problems studied in computational geometry and other areas including graph algorithms, geographical information systems (GIS) and robotics. Existing algorithms for most of the interesting shortest path problems are either very complex and/or have very large time and space complexities. Hence they are unappealing to practitioners and pose a challenge to theoreticians. This coupled with the fact that the geographic/spatial models are approximations of reality anyway and high-quality paths are favored over optimal paths that are "hard" to compute, approximation algorithms are suitable and necessary.

1.2 Related Work

Shortest path problems in computational geometry can be categorized by various factors which include the dimensionality of space, the type and number of objects or obstacles (e.g., polygonal obstacles, convex or non-convex polyhedra), and the distance measure used (e.g., Euclidean, number of links, or weighted distances). Several research articles, including surveys, have been written presenting the state-of-the-art in this active field. Here we discuss those contributions which relate more directly to our work; these are in particular stated in 2 and 3-dimensional weighted scenarios.

In two dimensions several variations of shortest path problems have been studied over the last two decades. This includes computing Euclidean shortest paths and answering shortest path queries between two points inside a simple polygon and amidst polygonal obstacles. Of particular interest is the *weighted region problem* (abbreviated as SPPS here) introduced by Mitchell and Papadrimitriou [21]; it is a natural generalization of the shortest path problem in polygonal domains. In their version [21] of the SPPS problem a planar triangulated subdivision is given consisting of n faces, where each face has a positive non-zero weight. Using an algorithm based on the "continuous Dijkstra's method" they [21] provide an approximation algorithm to compute a (weighted) ε-short path; it runs in $O(n^8 \log(nNW/w\epsilon))$ time using $O(n^4)$ space, where N is the largest integer coordinate of any vertex of the triangulation and W (w) is the maximum (minimum) weight of any face of the triangulation.

Lanthier et al. [17] have described several algorithms for the SPPS problem; the cost of the approximation is no more than the shortest path cost plus an (additive) factor of $W|LongestEdge|$, where $LongestEdge$ is the longest edge, W is the largest weight among all faces. As their experimental analysis shows these algorithms are also of practical value. They [19] also developed approximation algorithms for anisotropic shortest paths problems. Aleksandrov et al. [2] propose ε-approximation algorithms to solve the SPPS problem in $O((\frac{n}{\varepsilon^2} + \frac{n \log n}{\varepsilon}) \log \frac{1}{\varepsilon})$ time.

Given a set of pairwise disjoint polyhedra in $\Re^3$ and two points s and t, the problem of computing a shortest path between s and t that avoids the interiors of polyhedra is $NP - hard$ [5]. The shortest path problem amidst (disjoint) convex polyhedra can be solved in time exponential in the number of polyhedral objects as was shown by Sharir [23]. The NP-hardness and the large time complexities of 3-d shortest paths algorithms even for special problem instances have motivated the search for approximate solutions to shortest path problems.

Papadimitriou [22] was the first to study fully polynomial time approximation algorithms for the *Euclidean shortest path problem in three dimensions* (termed as ESP3D), where the interiors of the obstacles have infinite weight and the rest of the space has a uniform finite weight. Choi et al. [8; 9] present a refinement of the scheme proposed by [22] which provide an ε-approximation for ESP3D; their algorithm runs in $O((n^3 M \log M + (nM)^2).\mu(W))$ time, where $W = O(\log(\frac{n}{\varepsilon}) + L)$, $M = O(nL/\varepsilon)$ and each vertex of an obstacle is specified by an L-bit integer. Both of these algorithms divide the edges of the polyhedra into intervals by introducing Steiner points based on a geometric progression which depends on ε and the location of the source vertex, and then build a graph on these Steiner points. The ESP3D problem is solved by computing a shortest path in the graph using Dijkstra's algorithm [12]. Clarkson [7] pro-

vides a solution for the ESP3D problem. His algorithm produces an ε-short path and runs in $O(n^2\lambda(n)\log(n/\varepsilon)/(\varepsilon^4) + n^2\log n\rho\log(n\log\rho))$ time, where ρ is the ratio of the longest obstacle edge to the distance between the source and the target vertex, $\lambda(n) = \alpha(n)^{O(\alpha(n))^{O(1)}}$, $\alpha(n)$ is the inverse Ackermann's function.

Recently, there have been several results on approximation algorithms for computing Euclidean shortest paths on convex and nonconvex polyhedral surfaces. For example, Varadarajan and Agarwal [1] provide an algorithm that computes a path on a, possibly non-convex, polyhedron that is at most $7(1 + \varepsilon)$ times the shortest path length; it runs in $O(n^{5/3}\log^{5/3}n)$ time. They also present a slightly faster algorithm that returns a path which is at most $15(1 + \varepsilon)$ times the shortest path length. Algorithms of [17; 2] apply for Euclidean shortest path problems on polyhedral surfaces as well. We note that the Chen and Han's algorithm [6] can compute an exact Euclidean shortest path between two points on a polyhedral surface in $O(n^2)$ time, and recently Kapoor [15] claimed a faster algorithm.

1.3 Summary of our contributions

Our results improve upon previous results in a variety of ways:

1. *SPPS problem:* Considering the real-RAM model, we improve upon the result of Mitchell and Papadimitriou [21], by a factor of about n^7 in the problem size n. More specifically, our algorithm runs in $O(\frac{n}{\varepsilon}\log\frac{1}{\varepsilon}(\frac{1}{\sqrt{\varepsilon}} + \log n))$ time, whereas Mitchell and Papadimitriou's algorithm runs in $O(n^8\log(nNW/w\epsilon))$ time. Although the estimates include some geometric parameters, our result improves on that of [21] if $1/\varepsilon^{1.5}$ is less than $n^7/\log n$, which holds for most realistic problem instances. Although the improvement in the time complexity over [2] by a factor of $\frac{1}{\sqrt{\varepsilon}}$ may not look substantial, note that the margin for improvement is small, the constants arising out of geometric parameters are substantially better and the techniques developed in the process are novel and, as we show, are helpful in solving the 3-dimensional problem SP3D, which none of the previous techniques could.

2. *Unweighted SPPS problem:* We improve upon the results of Varadarajan and Agarwal [1] in the worst case time complexity as well as in terms of the approximation factors. They provide $7(1 + \varepsilon)$ and $15(1 + \varepsilon)$ approximations whereas ours is a true, i.e., $(1 + \varepsilon)$-approximation. For $\varepsilon > 1$ our algorithm runs in $O(n\log n)$ time, whereas theirs takes $O(n^{5/3}\log^{5/3}n)$ time.

3. *SP3D Problem:* We provide the first polynomial time ε-approximation scheme.

4. *Unweighted SP3D Problem:* For the unweighted version, namely the ESP3D problem, our worst case time complexities are better than those of [22; 8; 9] by at least a factor of n. Moreover, similar to the SPPS problem, our algorithm is simple and works for any pair of source and target vertex.

5. *Geometric parameters:* The constants that are hidden by the big-O notation may play an important role.

Our analysis (for SPPS) reveals fairly precisely the constants of the geometric portion of our algorithm (see e.g., Lemma 2.3). These constants depend on the geometric constellation of the given problem instance. The constants in previous approaches depend on ratios of maximum weight to minimum weight, maximum length v/s minimum length of edges, minimum angle of a face, etc. All our constants are averages over all faces and the dependence on the weights is square-root of the ratio of the maximum weight to the minimum weight. This in itself is significant, although it does not appear in the asymptotic bound.

6. *Simplicity:* Our algorithms are conceptually simple (in particular, compared to [21]), use few and standard geometric primitives and employ an readily implemented modification of Dijkstra's algorithm.

7. *Practicality potential:* Variants of the algorithm proposed here (for the SPPS), where the approximation factor is an additive factor, have been implemented and their practicality has been demonstrated (see [17]).

8. *Source independence:* Most of the previous approaches, including [22; 21; 1], build a graph over the surface, by using the knowledge of the location of the source vertex. This requires rebuilding of the graph if the location of the source vertex changes. Our approach does not require the knowledge of the location of the source or target vertex for building the graph, and hence one graph serves for any pair of source and target vertex.

To solve these problems, we employ the traditional technique of partitioning a continuous geometric search space into discrete combinatorial search space by designing an appropriate mesh. Simplifying, the general strategy taken here is to place Steiner points on edges and in the interior of faces (in 3-D), interconnect the Steiner points within each face and then show that an ε-short path exists approximating any (true) shortest path.

The Steiner point placement is novel (e.g., see the above discussion on source independence). One of the problems that arise is to place Steiner points near vertices and near edges in 3-D. As, near vertices (edges in 3-D), the distance between adjacent path vertices can become arbitrarily small, an infinite number of Steiner points would be required for the approximation (see also [16; 10]). Here, we address this problem by constructing "spheres" around the vertices and "spindles" around edges of the faces; ensuring that Steiner points are placed outside these regions. (Finding the right radius is important and non-trivial.) This allows us to put a lower bound on the length of the smallest possible edge that passes between two adjacent Steiner points and hence we are able to add a finite number of Steiner points. Now the challenge is to still show that an ε-approximation scheme is achievable.

Another issue that needs attention is that the constructed graph has considerable large size and it would be inappropriate to search it and hoping for the desired time complexities. We reduce the size of the search space during an execution of Dijkstra's algorithm by deriving geometric properties of Snell's law of refraction for a discrete domain. Snell's law is typically stated for continuous domains. We also describe a

variant of Dijkstra's algorithm in which the execution is restricted to a sparse set of potential edges, given that the preceding edge on a path is known. Employing both together and using geometric spanners we are able to construct a path of desired accuracy within the stated complexity bounds.

2. SPPS PROBLEM

2.1 Model

Let $\mathcal{P}$ be a polyhedral surface, whose faces are convex polygons. Let m be the number of edges and n be the number of faces of $\mathcal{P}$. Non-zero positive weights $w_1, \ldots, w_n$ are associated with faces $f_1, \ldots, f_n$ representing costs of traveling along them. The cost of traveling along an edge of $\mathcal{P}$, or otherwise the weight associated with an edge, is the minimum of the weights of the two neighboring faces. The cost of a path Π on $\mathcal{P}$ is defined by $\|\Pi\| = \sum_{i=1}^n w_i |\Pi_i|$, where $|\Pi_i|$ denotes the Euclidean length of the intersection $\Pi_i = \Pi \cap f_i$. (An edge is assumed to be part of the face from which it inherits its weight.)

Given two distinct points u and v on $\mathcal{P}$, a minimum cost path $\Pi(u, v)$ joining u and v is called *geodesic path*. In this setting, it is well known that any geodesic path is simple (non self-intersecting) and consists of a sequence of segments, whose endpoints are on the edges of $\mathcal{P}$. More precisely, each segment is of one of the following two types: 1) *face-crossing*: a segment which crosses a face joining two points on its boundary that lie on different edges 2) *edge-using*: a segment which (partially) follows an edge.

We define *linear paths* to be paths consisting of face-crossing and edge-using segments exclusively. A linear path $\Pi(u, v)$ is represented as a sequence of its segments $\{s_1, \ldots, s_{l+1}\}$ or/and as sequence of points $a_0, \ldots, a_{l+1}$ lying on adjacent edges of $\mathcal{P}$, that are endpoints of these segments, i.e., $s_i = (a_{i-1}, a_i)$, $u = a_0$, and $v = a_{l+1}$. Points a_i that are not vertices of $\mathcal{P}$ are called *bending points* of the path. The local behavior of a linear path around a bending point is described by two oriented angles, which we call *in-angle* and *out-angle* at the considered bending point.

Formally, let a be a bending point on a linear path Π lying on an edge e. Let s' be the segment preceding a on Π. Then we define the in-angle $\varphi = \varphi(a, \Pi)$ to be the acute clockwise angle between the normal to e at a and s'. Similarly, the out-angle ψ is the acute clockwise angle between the normal to e and the segment s'' succeeding a on Π. Note that in-angles and out-angles are always in $[-\pi/2, \pi/2]$. If the absolute value of an in-angle at some bending point a is $\pi/2$ then the segment preceding a is an edge-using segment. Similarly, if the absolute value of an out-angle is $\pi/2$ then the segment after a is an edge-using segment. In the case when Π is a geodesic path, angles φ and ψ are related as follows:

LEMMA 2.1. *(Snell's law) Let a be a bending point of a geodesic path Π that lies inside an edge e of $\mathcal{P}$. Let the segment preceding a in Π be s' and the one after a be s''. Let the costs of traveling along s' and s'' be w' and w'' respectively. Then*

$$w'' \sin \psi = w' \sin \varphi, \qquad (1)$$

where φ and ψ are the in-angle and out-angle at a.

In the next lemma we establish an estimate on the difference between out-angles of two given geodesic paths in terms of the difference between their in-angles. The result is somewhat counter intuitive as one may have expected that $|\varphi_1 - \varphi| \le \epsilon$ implies $|\psi_1 - \psi| \le \epsilon * constant$, for all angle pairs satisfying Snell's law and all ϵ, where the constant only depends on face weights[1].

LEMMA 2.2. *Let φ, ψ and φ_1, ψ_1 be two pairs of angles each satisfying (1). Assume that $w' > w''$. Then*

$$|\psi_1 - \psi| \le \frac{\pi w'}{2w''} \kappa(\varphi, \varphi_1) |\varphi_1 - \varphi|, \qquad (2)$$

where $\kappa(\varphi, \varphi_1) = \cos \frac{\varphi_1 + \varphi}{2} / \cos \frac{\psi_1 + \psi}{2}$.

Proof: Since each pair φ, ψ and φ_1, ψ_1 satisfies (1) we have

$$w''(\sin \psi_1 - \sin \psi) = w'(\sin \varphi_1 - \sin \varphi),$$

which is

$$w'' \sin \frac{\psi_1 - \psi}{2} \cos \frac{\psi_1 + \psi}{2} = w' \sin \frac{\varphi_1 - \varphi}{2} \cos \frac{\varphi_1 + \varphi}{2}. \quad (3)$$

Using the above relation and the inequalities

$$\frac{2}{\pi} \theta \le \sin \theta \le \theta,$$

that hold for any angle $\theta \in [0, \pi/2]$ we obtain the claim of the lemma

$$|\psi_1 - \psi| \le \pi |\sin \frac{\psi_1 - \psi}{2}|$$

$$= \frac{\pi w' \cos \frac{\varphi_1 + \varphi}{2}}{w'' \cos \frac{\psi_1 + \psi}{2}} |\sin \frac{\varphi_1 - \varphi}{2}|$$

$$\le \frac{\pi w'}{2w''} \kappa(\varphi, \varphi_1) |\varphi_1 - \varphi|.$$

$\square$

2.2 Discretization

DEFINITION 2.1. *Given a point $x \in \mathcal{P}$ the distance $d(x)$ of x to the set of edges of $\mathcal{P}$ is defined as the Euclidean distance $d(x)$ on $\mathcal{P}$ from x to $E \setminus E(x)$, where E is the set of the edges of $\mathcal{P}$ and $E(x)$ are the edges incident to x.*

For each vertex v of $\mathcal{P}$ we define a weighted radius

$$r(v) = \frac{w_{\min}(v)}{5w_{\max}(v)} d(v), \qquad (4)$$

where $w_{\max}(v)$ and $w_{\min}(v)$ are the maximum and minimum weights of the faces incident to v. By using the weighted radius $r(v)$ for each face incident to v we define a "small" isosceles triangle with two sides of length $\varepsilon r(v)$ incident to v. The small triangles around v form a star shaped polygon $S(v)$ which we call a *vertex-vicinity* of v.

Next, we describe the placement of Steiner points on edges, e, of $\mathcal{P}$. Assume that $e = (v', v'')$ and let $d(e) = \sup_{x \in e} d(x)$. Let M be a point on e, so that $d(M) = d(e)$. Such a point M can be easily computed in time proportional to the size of the two faces incident to e. The Steiner points on e are

---

[1] Mata and Mitchell's [20] algorithm is based on this, and their analysis as presented has gaps, since the angle between the cones after multiple refraction as claimed by them is incorrect

defined as follows. The point M is a Steiner point. On the edge-segment (v', M) we define points $p'_1, \ldots, p'_{k'}$ so that $|v'p'_1| = \varepsilon r(v')$, and $|p'_{i-1}p'_i| = \varepsilon d(p'_{i-1})$, for $i = 2, \ldots k'$. Note that the first Steiner point p'_1 coincides with the vertex of the vertex-vicinity $S(v')$ that lies on e. The Steiner points $p''_1, \ldots, p''_{k''}$ on the edge-segment (v'', M) are defined in an analogous manner. To simplify notation, we denote Steiner points on e following their order from v' to v'' by $p_1, \ldots, p_k$, where $k = k' + k'' + 1$ and $M = p_{k'+1}$. Using the above procedure we insert Steiner points on all edges of $\mathcal{P}$. In the next lemma we estimate the number of Steiner points inserted on an edge of $\mathcal{P}$ and consequently the total number of Steiner points on $\mathcal{P}$. We denote by $E(\mathcal{P})$ the set of edges of $\mathcal{P}$.

LEMMA 2.3. *(a) The number of Steiner points inserted on an edge $e = (v', v'')$ does not exceed $3 + C(e)\frac{1}{\varepsilon}\log_2\frac{2}{\varepsilon}$, where the constant $C(e) < 4(|e|/d(e))\log_2(|e|/\sqrt{r(v')r(v'')})$.*
(b) The total number of the Steiner points inserted on $\mathcal{P}$ is bounded by $\tilde{C}(\mathcal{P})\frac{m}{\varepsilon}\log_2\frac{2}{\varepsilon}$, where $\tilde{C}(\mathcal{P}) - 3$ does not exceed the average of the constants $C(e)$ on $E(\mathcal{P})$, i.e., $\tilde{C}(\mathcal{P}) \leq 3 + \frac{1}{m}\sum_{e=(v',v'')\in E(\mathcal{P})} C(e)$.

Proof: Recall that we denoted $d(e) = \sup_{x \in e} d(x)$ and let $\tau(e)$ be the ratio $d(e)/|e|$. Note that $\tau(e)$ is a positive number smaller than $1/2$ and in some sense measures the "thickness" of the two faces around e. Let $M' \in e$ be the closest to v' point on e, such that $d(M') = d(e)$. According to our construction the point M' is between v' and M or coincides with M. Let $p'_1, \ldots, p'_{k'_1}$ be the Steiner points on the subsegment (v', M') and $p'_{k'_1+1}, \ldots, p'_{k'}$ are those in (M', M). Then by the fact that $d(x)/|v'x|$ is a decreasing function in (v', M') we have

$$|v'p'_i| \geq (1 + \varepsilon\tau(e))|v'p'_{i-1}|,$$

and therefore

$$|v'p'_i| \geq \varepsilon r(v')(1 + \varepsilon\tau(e))^{i-1} \quad \text{for} \quad i = 1, \ldots, k'_1 + 1.$$

Since $|v'p'_{k'_1+1}| \leq |v'M|$ we have

$$k'_1 \leq \log_{(1+\varepsilon\tau(e))}(|v'M|/\varepsilon r(v')). \tag{5}$$

By the definition the points $p'_{k'_1+1}, \ldots, p'_{k'}$ are equidistantly placed in (M', M) at a distance $\varepsilon d(e)$ and therefore

$$k' - k'_1 - 1 \leq \frac{|M'M|}{\varepsilon d(e)}.$$

Summing these two inequalities we obtain

$$k' - 1 \leq \log_{(1+\varepsilon\tau(e))}(|v'M|/\varepsilon r(v')) + \frac{|M'M|}{\varepsilon d(e)}. \tag{6}$$

To estimate the number $k = k' + k'' + 1$ of Steiner points inserted in e we sum (6) with the analogous inequality for k'', use that $|M'M''| \leq |e|$ and $4|v'M'||v''M''| \leq |e|^2$ and obtain

$$
\begin{aligned}
k - 3 & \\
\leq \quad & \log_{(1+\varepsilon\tau(e))}(|v'M|/\varepsilon r(v')) + \frac{|M'M|}{\varepsilon d(e)} \\
& + \log_{(1+\varepsilon\tau(e))}(|v''M|/\varepsilon r(v'')) + \frac{|M''M|}{\varepsilon d(e)} \\
\leq \quad & \log_{(1+\varepsilon\tau(e))} \frac{|v'M||v''M|}{\varepsilon^2 r(v')r(v'')} + \frac{|M'M''|}{\varepsilon d(e)} \\
\leq \quad & 2\log_{(1+\varepsilon\tau(e))} \frac{|e|}{2\varepsilon\sqrt{r(v')r(v'')}} + \frac{1}{\varepsilon\tau(e)}.
\end{aligned}
$$

Then the estimate stated in the lemma is derived from this inequality using properties of the logarithm function and the fact that $\varepsilon\tau(e) \leq 1/2$. Namely

$$
\begin{aligned}
k - 3 &\leq 2\log_{(1+\varepsilon\tau(e))} \frac{|e|}{2\varepsilon\sqrt{r(v')r(v'')}} + \frac{1}{\varepsilon\tau(e)} \\
&\leq \frac{1}{\varepsilon\tau(e)}(1 + \frac{5}{2}\ln\frac{|e|}{2\varepsilon\sqrt{r(v')r(v'')}}) \leq C(e)\frac{1}{\varepsilon}\log_2\frac{2}{\varepsilon},
\end{aligned}
$$
$$\tag{7}$$

where $C(e) < 4(|e|/d(e))\log_2(|e|/\sqrt{r(v')r(v'')})$. The estimate on the total number of Steiner points on $\mathcal{P}$ is obtained by summing (7) for all edges of $\mathcal{P}$. $\quad\square$

The set of Steiner points defines a set of *Steiner intervals*. The essential property of the Steiner point set on $\mathcal{P}$ is that any face-crossing segment s having an endpoint in a Steiner interval (p_i, p_{i+1}), $1 \leq i \leq k - 1$, satisfies the inequality

$$|p_ip_{i+1}| \leq \varepsilon|s|. \tag{8}$$

The above property provides an upper bound on the maximum angle θ at which a Steiner interval (p_i, p_{i+1}) is seen from a point on the boundary of the quadrilateral formed by the two triangles neighboring that interval.

LEMMA 2.4. *Let (p_i, p_{i+1}) be a Steiner interval on an edge e and x be a point on the boundary of the union of the two faces neighboring e. Then*

$$\angle p_i x p_{i+1} \leq \frac{\pi}{2}\varepsilon. \tag{9}$$

Next we show that the density of the Steiner points is sufficient to approximate face-crossing segments with segments joining pairs of Steiner points with ε accuracy. More precisely, for a face-crossing segment (a, b) that does not intersect vertex vicinities we show that:

LEMMA 2.5. *Let (a, b) be a face-crossing segment that does not belong to a vertex-vicinity.*
a) *If (a, b) joins a pair of Steiner intervals say (p_i, p_{i+1}) and (q_j, q_{j+1}) then*

$$\max\left[\min(|p_iq_j|, |p_iq_{j+1}|), \min(|p_{i+1}q_j|, |p_{i+1}q_{j+1}|)\right]$$
$$\leq (1 + \varepsilon)|ab|. \tag{10}$$

b) *If (a, b) joins a segment (v, p_1) between a vertex v and a Steiner point p_1 with a Steiner interval (q_j, q_{j+1}) then*

$$\max\left[(|p_1q_j|, |p_1q_{j+1}|)\right] \leq (1 + \varepsilon)|ab| + \varepsilon r(v) \tag{11}$$

c) *If (a, b) joins two segments (v, p_1) and (v', q_1), where v and v' are vertices and p_1, q_1 are their neighbour Steiner points then*

$$|p_1q_1| \leq |ab| + \varepsilon(r(v) + r(v')) \tag{12}$$

2.3 Discrete paths

In Section 2.2 we have described a particular method for placing Steiner points on edges of $\mathcal{P}$. This suggests the definition of a graph $G_\varepsilon = (V(G_\varepsilon), E(G_\varepsilon))$ whose vertex set $V(G_\varepsilon)$ includes all Steiner points plus the vertices of $\mathcal{P}$. The set of edges of $E(G_\varepsilon)$ consists of all face-crossing segments joining pairs of Steiner points plus edge-using segments joining either a pair of neighboring Steiner points or a vertex of $\mathcal{P}$ with its neighboring Steiner point. Now any path in G_ε is referred to as a *discrete path* on $\mathcal{P}$. Note that discrete paths

are linear as well.

An ε-*short* path for a path $\tilde{\Pi}$ on $\mathcal{P}$ is a discrete path whose cost is at most $(1+\varepsilon)\|\tilde{\Pi}\|$. In this section we first show the existence of a 3ε-short discrete path for any linear path on $\mathcal{P}$. Then, by using a spanner technique, we define a subgraph G_ε^* of G_ε that has a smaller set of edges but still contains 7ε-short discrete path for any linear path. The graph G_ε^* will be used in the next section where we present a pruned Dijkstra algorithm for approximating geodesic paths on $\mathcal{P}$.

Next we introduce some notions about the structure of linear paths as necessary for our considerations. Consider a linear path $\tilde{\Pi}(v_0, v)$ joining two different vertices of $\mathcal{P}$. This path crosses vertex-vicinities $S(v_0)$ and $S(v)$ and possibly a number of other vertex-vicinities. The path $\tilde{\Pi}$ can be partitioned into portions $\tilde{\Pi}(S(v_i))$, lying fully inside *vertex-vicinities*, $S(v_i)$, and those, called *between-vertex-vicinities* $\tilde{\Pi}(S(v_i), S(v_{i+1}))$, connecting vertex-vicinities $S(v_i)$ and $S(v_i + 1)$. [2] Assume now that the path $\tilde{\Pi}(v_0, v)$ consists of $k + 2$ vertex-vicinity and $k + 1$ between-vertex-vicinity portions.

Next, we introduce the notion of a discrete path that *neighbours* a given linear path $\tilde{\Pi}(v_0, v)$. We describe first a discrete path that neighbours a between-vertex-vicinities portion. Consider a between-vertex-vicinity portion $\tilde{\Pi}(S(v_i), S(v_{i+1}))$ for some $0 \le i \le k$. This portion consists of a sequence of face-crossing and edge-using segments that do not intersect vertex vicinities. Each of these segments joins two Steiner intervals and thereafter the portion $\tilde{\Pi}(S(v_i), S(v_{i+1}))$ defines a sequence of Steiner intervals. A discrete path $\Pi(S(v_i), S(v_{i+1}))$ neighbours the between-vertex-vicinities linear path $\tilde{\Pi}(S(v_i), S(v_{i+1}))$ if it joins a Steiner point on the boundary of the vertex-vicinity $S(v_i)$ with a Steiner point on the boundary of $S(v_{i+1})$ and crosses the same sequence of Steiner intervals. Informally, when the linear path crosses a Steiner interval its discrete neighbour snaps to one of the Steiner points forming that interval. In general, we use the following definition.

DEFINITION 2.2. *A discrete path $\Pi(v_0, v)$ is a neighbour of the linear path $\tilde{\Pi}(v_0, v)$ given as above, if its between-vertex-vicinities portions neighbour corresponding between-vertex-vicinities portions of $\tilde{\Pi}$ and its vertex-vicinity portions are two edge paths joining between-vertex-vicinities portions through the corresponding vertex.*

Now we state and prove (see the full version [3] also for a formal definition of path neighbour) the central theorem of this section.

THEOREM 2.1. *Let $\tilde{\Pi}(v_0, v)$ be a linear path joining two different vertices on $\mathcal{P}$. There exists a discrete 3ε-short path $\Pi(v_0, v)$ that neighbours $\tilde{\Pi}$.*

Theorem 2.1 provides a method for approximating geodesic paths on $\mathcal{P}$. We may consider the graph $G_\varepsilon = (V(G_\varepsilon), E(G_\varepsilon))$

---

[2]Note that according to our definition a vertex-vicinity portion $\tilde{\Pi}(S(v))$ can be empty although the path crosses $S(v)$ as well it can contain segments outside $S(v)$ that lie between the first and the last intersection point of $\tilde{\Pi}$ and $S(v)$.

and assign weights on its edges equal to the cost of traveling along them. Then by Theorem 2.1 the cost of a shortest path in G_ε between any two of its nodes does not exceed $(1 + 3\varepsilon)$ times the cost of the geodesic path on $\mathcal{P}$. Then Dijkstra's algorithm can be used for finding shortest paths in G_ε in time $O(|E(G_\varepsilon)| + |V(G_\varepsilon)| \log |V(G_\varepsilon)|)$. Therefore, we may approximate geodesic paths on $\mathcal{P}$ in time $O(\frac{n}{\varepsilon} \log \frac{1}{\varepsilon}(\frac{1}{\varepsilon} \log \frac{1}{\varepsilon} + \log(\frac{n}{\varepsilon} \log \frac{1}{\varepsilon})))$.

COROLLARY 2.1. *For $\varepsilon = \frac{1}{3}$ we obtain an $O(n \log n)$ algorithm that finds a path of cost at most twice that of the shortest.*

A weighted shortest path can have $O(\Omega(n^2))$ segments. This does however not contradict our time bounds as in such a case many segments will cross the same Steiner interval and are thus replaced by a single approximating segment. For the remainder we will assume that $\varepsilon < 1$. Note however that our approach so far does not use any properties of geodesic paths on $\mathcal{P}$. It turns out that using the additional geometric information of geodesic paths (provided by Lemma 2.1) we are able to prove a useful characterization of the paths in G_ε that neighbour geodesic paths on $\mathcal{P}$ and thereafter direct our search. This will result in an $O(\frac{n}{\varepsilon} \log \frac{1}{\varepsilon}(\frac{1}{\sqrt{\varepsilon}} + \log n))$ time approximation algorithm.

DEFINITION 2.3. *A discrete path $\Pi(v_0, v)$ is a discrete geodesic path if it neighbours a geodesic path joining v_0 and v.*

Theorem 2.1 directly implies the following corollary.

COROLLARY 2.2. *For any geodesic path $\tilde{\Pi}(v_0, v)$ there exist a discrete geodesic path $\Pi(v_0, v)$ such that $\|\Pi\| \le (1 + 3\varepsilon)\|\tilde{\Pi}\|$.*

Corollary 2.2 suggests that in order to approximate geodesic paths on $\mathcal{P}$ it suffices to search in the class of discrete geodesic paths. So we need an efficient procedure to compute and search that class. First we show that discrete geodesic paths satisfy a constraint similar to that established for geodesic paths in Lemma 2.1. More precisely given a discrete geodesic path Π and a segment (a, b) of Π, we can prove that there is a cone $C(a, b)$ with apex b and angle depending on ε and weights around b, that contains the next bending point c of Π. Note that in the case of geodesic paths any of its segments completely determines the direction of its next segment, i.e., the angle of the corresponding cone is zero.

Let $\Pi(v_0, v)$ be a discrete geodesic path. By definition there is a geodesic path neighboring Π. Let a, b, and c be three consecutive bending points of Π, where b does not neighbour a vertex of $\mathcal{P}$. Here a and c are Steiner points (not vertices of $\mathcal{P}$). According to Lemma 2.1, the segment (a, b) defines a unique direction at which a geodesic path containing (a, b) must continue after b. Except for the case where that direction is along the edge containing b, a unique face-crossing segment (b, c_0) is defined so that the in-angle φ and the out-angle ψ at b of the sub-path $\{a, b, c_0\}$ satisfy Lemma 2.1. The next lemma establishes an estimate on the size of the angle $\angle cbc_0$.

LEMMA 2.6. *Let $\varepsilon < 1$ and a, b and c be three consecutive bending points on a discrete geodesic path Π, where the Steiner point b does not neighbour a vertex of $\mathcal{P}$. Let c_0 be*

a point on the geodesic direction determined by the segment (a, b) (on the ray starting at b and with an out-angle ψ that satisfies Lemma 2.1). Then

$$\angle cbc_0 \leq \pi(1 + \sqrt{\frac{\pi w'}{2w''}})\sqrt{\varepsilon}, \tag{13}$$

w', w'' are the weights of the faces incident to b.

Proof: Let us denote by a_1, a_2, b_1, b_2, and c_1, c_2 the left and right Steiner points neighboring a, b and c respectively. By the definition Π neighbors a geodesic path $\tilde{\Pi}$ and let the bending points of $\tilde{\Pi}$ sharing common Steiner intervals with a, b and c be $\tilde{a}$, $\tilde{b}$ and $\tilde{c}$. Thereby we have $\tilde{a} \in (a_1, a_2)$, $\tilde{b} \in (b_1, b_2)$, and $\tilde{c} \in (c_1, c_2)$.
The segment (a, b) is either a face-crossing or an edge-using segment. We consider these two possibilities separately.

Case 1. Let (a, b) be a face-crossing segment.
First we shall estimate the differences between in-angles of the segments (a, b) and $(\tilde{a}, \tilde{b})$. Let us denote these two in angles by φ and $\tilde{\varphi}$. Recall that φ and $\tilde{\varphi}$ are clockwise angles between segments (a, b) and $(\tilde{a}, \tilde{b})$ and normals at b and $\tilde{b}$. We are going to show that

$$|\varphi - \tilde{\varphi}| \leq \pi \varepsilon. \tag{14}$$

One easily observes that $|\varphi - \tilde{\varphi}|$ always equals the angle θ between lines containing (a, b) and $(\tilde{a}, \tilde{b})$ and so it is enough to show $\theta \leq \pi \varepsilon$. Let the intersection point between lines containing (a, b) and $(\tilde{a}, \tilde{b})$ by x (if these lines do not intersect θ is zero and the claim is trivial). Then

$$\theta = \angle ax\tilde{a} \leq \angle ab\tilde{a} + \angle \tilde{b}\tilde{a}b,$$

where we have equality if the point $x = (a, b) \cap (\tilde{a}, \tilde{b})$ and inequality if (a, b) and $(\tilde{a}, \tilde{b})$ do not intersect. Now it follows that $\theta \leq \pi \varepsilon$, since by Lemma 2.4 both $\angle ab\tilde{a}$ and $\angle \tilde{b}\tilde{a}b$ are at most $\pi \varepsilon / 2$.
The same argument implies that

$$|\psi_c - \tilde{\psi}| \leq \pi \varepsilon, \tag{15}$$

where ψ_c and $\tilde{\psi}$ are the out-angles of the segments (b, c) and $(\tilde{b}, \tilde{c})$.
Using that the angle $\angle cbc_0$ is equal to $|\psi - \psi_c|$ triangle inequality and (15) we obtain

$$\angle cbc_0 = |\psi - \psi_c| \leq |\psi_c - \tilde{\psi}| + |\psi - \tilde{\psi}| \leq \pi \varepsilon + |\psi - \tilde{\psi}|. \tag{16}$$

Next we estimate $|\psi - \tilde{\psi}|$. Lemma 2.2 and (14) imply

$$|\psi - \tilde{\psi}| \leq \frac{\pi w'}{2w''} \kappa(\varphi, \tilde{\varphi}) \pi \varepsilon$$

and by using a simple inequality $\kappa(\varphi, \tilde{\varphi}) \leq \pi / (\pi - |\psi + \tilde{\psi}|)$ we have

$$|\psi - \tilde{\psi}| \leq \frac{\pi^3 w' \varepsilon}{2w''(\pi - |\psi + \tilde{\psi}|)}. \tag{17}$$

On the other hand, since $\psi, \tilde{\psi} \in [-\pi/2, \pi/2]$ we have $|\psi - \tilde{\psi}| + |\psi + \tilde{\psi}| \leq \pi$ and therefore

$$|\psi - \tilde{\psi}| \leq \min\left(\pi - |\psi + \tilde{\psi}|, \frac{\pi^3 w' \varepsilon}{2w''(\pi - |\psi + \tilde{\psi}|)}\right)$$
$$\leq \pi\sqrt{\frac{\pi w' \varepsilon}{2w''}}. \tag{18}$$

We obtain the lemma in the case where (a, b) is a face-crossing segment by substitution of (18) in (16).

Case 2. Let (a, b) be an edge-using segment.
In this case by Lemma 2.1 we know that the out-angles ψ and $\tilde{\psi}$ are equal. Thus

$$\angle cbc_0 = |\psi - \psi_c| = |\tilde{\psi} - \psi_c|.$$

But as we have proved $|\tilde{\psi} - \psi_c| < \pi \varepsilon$ and the lemma follows.
$\square$
The result of Lemma 2.6 can be summarized in the following way. If Π is a discrete geodesic path and (a, b) is a segment of Π, so that b does not neighbour a vertex of $\mathcal{P}$ then the cone defined using (13) contains the next segment (b, c) of Π. Hereafter, we call the set of edges (b, c) that satisfy (13) with respect to (a, b) *geodesic cone* and denote it by $GCone(a, b)$. Note that if b neighbours or is a vertex $\mathcal{P}$ then the lemma provides no information and we assume, that $GCone(a, b)$ consists of all edges incident to b except (a, b).

Let a_0 be any fixed node of G_ε. Based on Lemma 2.6 we define a class of paths $\mathcal{C}(a_0)$ consisting of all discrete paths that originate from a_0 and satisfy Lemma 2.6. Clearly, $\mathcal{C}(a_0)$ contains discrete geodesic paths originating from a_0. An inductive description of $\mathcal{C}(a_0)$ is as follows:

DEFINITION 2.4. *The class of paths $\mathcal{C}(a_0)$ is defined by:*

1. *Edges of G_ε incident to a_0 belong to $\mathcal{C}(a_0)$.*

2. *A discrete path $\Pi' = \{a_0, a_1, \ldots, a_{k-1}, a_k\}$, $k > 1$ belongs to $\mathcal{C}(a_0)$ iff $\Pi = \{a_0, a_1, \ldots, a_{k-1}\} \in \mathcal{C}(a_0)$ and $(a_{k-1}, a_k) \in GCone(a_{k-2}, a_{k-1})$.*

2.4 A Pruned Dijkstra Algorithm

We have shown that geodesic paths originating from a fixed node a_0 of G_ε to all other nodes can be approximated by finding single source shortest paths in $\mathcal{C}(a_0)$. To compute those, we employ a *pruned Dijkstra* algorithm whose search is restricted to the class $\mathcal{C}(a_0)$. Note that the inductive definition of $\mathcal{C}(a_0)$ perfectly meets that purpose. However, we cannot prove the claimed efficiency bound since the number of the edges in cones $Cone(a, b)$ still can be large. To obtain an efficient algorithm we define and employ a class of paths $\mathcal{C}^*(a_0)$ that contains an ε-approximation path for each path in $\mathcal{C}(a_0)$. Additionally the class $\mathcal{C}^*(a_0)$ will be defined so that the pruned Dijkstra algorithm will perform fewer updates per step while searching it. The definition of $\mathcal{C}^*(a_0)$ resembles Definition 2.4, but here the sets of edges, $Cone^*(\Pi)$, that can continue a given path $\Pi \in \mathcal{C}^*(a_0)$ are subsets of the cones from Definition 2.4 of size $O(\varepsilon^{-\frac{1}{2}})$. We specify sets $Cone^*(\Pi)$ in detail after the definition.

DEFINITION 2.5. *The class $\mathcal{C}^*(a_0)$ consists of:*

1. *Edges of G_ε incident to a_0 are one edge paths in $\mathcal{C}^*(a_0)$.*

2. *A discrete path $\Pi' = \{a_0, a_1, \ldots, a_{k-1}, a_k\}$, $k > 1$ belongs to $\mathcal{C}^*(a_0)$ iff*

$$\Pi = \{a_0, a_1, \ldots, a_{k-1}\} \in \mathcal{C}^*(a_0) \quad \text{and}$$
$$(a_{k-1}, a_k) \in Cone^*(\Pi). \tag{19}$$

Next, given a path $\Pi = \{a_0, a_1, \ldots, a_{k-1}, \}$, $k > 1$ in $\mathcal{C}^*(a_0)$ we define the set $Cone^*(\Pi) = Cone^*(\Pi, a_{k-1})$. Intuitively,

it consists of an ε-spanner of some of geodesic cones $Cone(a_{k-i}, a_{k-1})$, $i > 1$. Let us introduce first the notion of an ε-spanner of a cone of edges. A cone $Cone(b)$ with apex b is a set of all edges incident to b that belong to the intersection of one of the faces around b. Equivalently, $Cone(b)$ is a sequence of consecutive edges around b lying in one of the faces that neighbour b. The angle between the first and the last edge in the sequence $Cone(b)$ is called the angle of $Cone(b)$. Note that by our definition the angle of a cone can not exceed π.

An ε-spanner $Cone_\varepsilon(b)$ of $Cone(b)$ with an angle $\theta \in (0, \pi)$ is defined as a subset of at most $\lceil 2\theta/\pi\varepsilon \rceil$ edges of $Cone(b)$ as follows. The $Cone(b)$ is partitioned into $j \leq \lceil 2\theta/\pi\varepsilon \rceil$ sub-cones $Cone_1(b), \ldots, Cone_j(b)$, so that each of these cones contains at least one edge and their angles do not exceed $\pi\varepsilon/2$. Such a partition is always possible as a result of Lemma 2.4. Then $Cone_\varepsilon(b)$ is defined as the set of the shortest length edges $s_1, \ldots, s_j$, in the sub-cones $Cone_1(b), \ldots, Cone_j(b)$. The lemma below states an important property of ε-spanners.

LEMMA 2.7. *Let $Cone(b)$ and $Cone_\varepsilon(b)$ be a cone of edges and its ε-spanner. For any edge (b, c) in $Cone(b)$ there is an edge (b, c') in $Cone_\varepsilon(b)$ so that*

$$|bc'| \leq |bc| \leq |bc'| + |cc'| \leq (1 + \frac{\pi\varepsilon}{2})|bc|. \qquad (20)$$

The following cases arise for the last node and edge of the path $\Pi = \{a_0, a_1, \ldots, a_{k-1}, \}$:

Case 1: The node a_{k-1} is a vertex of $\mathcal{P}$.
In this case the set $Cone^*(\Pi, a_{k-1})$ consists of all edges incident to a_{k-1} except (a_{k-2}, a_{k-1}). The number of edges in $Cone^*(\Pi, a_{k-1})$ in this case is equal to the degree of the vertex a_{k-1} in $\mathcal{P}$.

Case 2: The node a_{k-1} neighbours a vertex of $\mathcal{P}$.
In this case the edges incident to a_{k-1} can be partitioned into two cones – one per each of the two faces incident to a_{k-1}. Then the set $Cone^*(\Pi, a_{k-1})$ is defined as the union of the ε-spanners of these cones. The number of edges in $Cone^*(\Pi, a_{k-1})$ in this case does not exceed $\lceil 4/\varepsilon \rceil$.

Case 3: The node a_{k-1} neither neighbours nor is a vertex of $\mathcal{P}$.
In this case the definition of $Cone^*(\Pi, a_{k-1})$ depends on the type of the segment (a_{k-2}, a_{k-1}) and we consider the two possible sub-cases.

Case 3.1: The edge (a_{k-2}, a_{k-1}) is an edge-using segment. In this case the set $Cone^*(\Pi, a_{k-1})$ depends on the first node a_{k-i}, $i > 2$ backward on the path Π, for which (a_{k-i}, a_{k-1}) is a face-crossing segment (see Figure 6). The set $Cone^*(\Pi, a_{k-1})$ consists of the edge joining a_{k-1} with its neighboring Steiner point different of a_{k-2} plus ε-spanners of the geodesic cones of the segments (a_{k-i}, a_{k-1}) and (a_{k-i+1}, a_{k-1}). Note that in some of the cases the corresponding geodesic cones are empty and so are their spanners.

To estimate the size of the set $Cone^*(\Pi, a_{k-1})$ in this case we use the estimates for the angles of the geodesic cones obtained in Lemma 2.6. Then we have estimates on the size of the spanners based on the size of the cone

angles. We add these estimates, use that $\varepsilon < 1$ and obtain

$$|Cone^*(\Pi, a_{k-1})| \leq 2 + \pi^2(2 + \sqrt{\frac{\pi w'}{2w''}})\varepsilon^{-\frac{1}{2}}. \qquad (21)$$

Here w' is the weight of the face containing (a_{k-i}, a_{k-1}) and w'' is the weight of the other face incident to a_{k-1}.

Case 3.2: The edge (a_{k-2}, a_{k-1}) is a face-crossing segment. In this case the set $Cone^*(\Pi, a_{k-1})$ consists of the two edges joining a_{k-1} with its neighboring Steiner points plus the ε-spanner of the geodesic cone $GCone(a_{k-2}, a_{k-1})$. The angle of the geodesic cone was estimated in Lemma 2.6 and according to the construction of ε-spanner and (13) the size of $Cone^*(\Pi, a_{k-1})$ is at most

$$|Cone^*(\Pi, a_{k-1})| \leq 3 + \pi^2(1 + \sqrt{\frac{\pi w'}{2w''}})\varepsilon^{-\frac{1}{2}}, \qquad (22)$$

where w' is the weight of the face containing (a_{k-2}, a_{k-1}) and w'' is the weight of the other face incident to a_{k-1}.

In the next lemma we establish that paths in the class $\mathcal{C}(a_0)$ are approximated by paths in $\mathcal{C}^*(a_0)$.

LEMMA 2.8. *Let Π be a discrete geodesic path starting at a_0 and consisting of k edges $\Pi = \{a_0, a_1, \ldots, a_k\}$. The class $\mathcal{C}^*(a_0)$ contains a path Π^* of the form $\Pi^* = \{a_0, a_1, \ldots, a_2, \ldots, a_{k-1}, \ldots, a_k\}$, where each of the sub-paths $\Pi^*(a_{i-1}, a_i)$, for $i = 1, \ldots, k$, consists either of a single edge (a_{i-1}, a_i) or of a single face-crossing edge (a_{i-1}, a_i') plus a sequence of edges joining consecutive Steiner points between a_i' and a_i along an edge of $\mathcal{P}$. The cost of the path Π^* satisfies*

$$\|\Pi^*\| \leq (1 + \pi\varepsilon/2)\|\Pi\|. \qquad (23)$$

Next we present an efficient algorithm for finding shortest paths in the class $\mathcal{C}^*(a_0)$ from a_0 to all other nodes of G_ε. Actually, this is Dijkstra's single source shortest paths algorithm restricted to search in the class $\mathcal{C}^*(a_0)$ only. The algorithm takes as input the surface $\mathcal{P}$, the set of Steiner points V_ε as defined in Section 2.2 and a fixed Steiner point a_0. It outputs a tree $SPT(\mathcal{C}^*(a_0))$ routed at a_0 and consisting of single source shortest paths in the class $\mathcal{C}^*(a_0)$. The only difference between the algorithm and the classical Dijkstra's algorithm is that this one does not update all possible continuations of the latest output shortest path, but only the nodes forming paths in $\mathcal{C}^*(a_0)$ using edges $(a, b) \in Cone^*(\Pi(a), a)$. As the classical Dijkstra's algorithm, our modified algorithm, called PrunedDjikstra, employs a priority queue Q.

We may implement Q (see, [13]) so that the amortized time for each update in Step 2.2 is a constant and therefore it is proportional to the size of the corresponding set $Cone^*(\Pi(a), a)$ if a was extracted from the queue. According to the definition of the sets $Cone^*(\Pi(a), a)$ their sizes are:

1. $|Cone^*(\Pi(a), a)| = O(\varepsilon^{-\frac{1}{2}})$ if a is a node that does not neighbour nor is a vertex of $\mathcal{P}$.

2. $|Cone^*(\Pi(a), a)| = O(\varepsilon^{-1})$ if a neighbours a vertex of $\mathcal{P}$.

3. $|Cone^*(\Pi(a), a)| = \deg(a)$ if a is a vertex of $\mathcal{P}$, where $\deg(a)$ denotes the number of edges incident to a in $\mathcal{P}$.

Therefore the total time for the implementation is proportional to the $\sum_{a \in V_\varepsilon} |Cone^*(\Pi(a), a)|$ which equals $O(\frac{n}{\varepsilon^{\frac{3}{2}}} \log \frac{1}{\varepsilon})$.

LEMMA 2.9. *Algorithm PrunedDjikstra outputs a rooted spanning tree $SPT(\mathcal{C}^*(a_0))$ of G_ε and for each $a \in V_\varepsilon$, the cost of the shortest path from a_0 to a in the class $\mathcal{C}^*(a_0)$. The algorithm runs in $O(\frac{n}{\varepsilon} \log \frac{1}{\varepsilon}(\frac{1}{\sqrt{\varepsilon}} + \log n))$ time.*

In summary, given a polyhedral surface $\mathcal{P}$ as defined in Section 2.1 and an approximation parameter $\varepsilon > 0$ we discretize $\mathcal{P}$ by adding a set of Steiner points as described in Section 2.2. This results in an augmented vertex set V_ε. Then we consider a fixed source point $a_0 \in V_\varepsilon$ and let $Cost(a; \mathcal{P})$ denote the cost of the shortest path from a_0 to $a \in V_\varepsilon$ on $\mathcal{P}$. To find approximate shortest paths we simply run Algorithm PrunedDjikstra above and report the paths $\Pi^*(a_0, a)$ from the spanning tree $SPT(\mathcal{C}^*(a_0))$ computed by the algorithm. According to Lemma 2.9, Lemma 2.8 and Theorem 2.1 we have

$$
\begin{aligned}
\|\Pi^*(a_0, a)\| &= \min\{\|\Pi\| : \Pi \in \mathcal{C}^*(a_0)\} \\
&\leq (1 + \pi\varepsilon/2)\min\{\|\Pi\| : \Pi \in \mathcal{C}(a_0)\} \\
&\leq (1 + 3\varepsilon)(1 + \pi\varepsilon/2)Cost(a; \mathcal{P}) \\
&\leq (1 + 15\varepsilon)Cost(a; \mathcal{P})
\end{aligned}
$$

Thus we have established the following theorem:

THEOREM 2.2. *Let $\mathcal{P}$ be a weighted polyhedral surface and $0 < \varepsilon \leq 1$. The single source shortest paths on $\mathcal{P}$ can be approximated with accuracy factor $(1 + 15\varepsilon)$ in $O(\frac{n}{\varepsilon} \log \frac{1}{\varepsilon}(\frac{1}{\sqrt{\varepsilon}} + \log n))$ time.*

3. WRP-3D PROBLEM

3.1 Introduction
This section presents an approximation scheme for a given triangulation (tetrahedralization) $\mathcal{D}$ in 3-dimensional Euclidean space $\Re^3$ consisting of edges, faces and tetrahedra. The domain is discretized by placing a set of Steiner points on the faces and edges of the tetrahedra and inserting a set of segments across tetrahedra that interconnects pairs of Steiner points. The placement and the number of Steiner points depends upon the desired approximation factor ε as well as the geometry and weights of the given configuration. Following the approach taken in solving the SPPS problem, to solve WRP-3D we need to address a) the placement of Steiner points, b) the discretization of Snell's law in 3-dimensional space, and c) the approximation quality of paths computed by the Pruned Dijkstra's Algorithm. Due to the lack of space, we can only sketch the main ideas here and refer the reader to the full version [3].

3.2 Preliminaries
Let $\mathcal{D}$ be a connected triangulation of space in three dimensions, with m edges, f faces (triangles) and n tetrahedra $t_1, \ldots, t_n$. Positive weights $w_1, \ldots, w_n$ are associated with tetrahedra representing the costs of traveling inside them. The cost of traveling along a face or an edge of a tetrahedron is the minimum of the weights of the neighboring tetrahedra. The cost of a path π on $\mathcal{D}$ is defined by $\|\pi\| = \sum_{i=1}^{n} w_i|\pi_i|$, where $|\pi_i|$ denotes the Euclidean length of the intersection $\pi_i \cap t_i$. Given two vertices u and v in $\mathcal{D}$ the geodesic path $\pi(u, v)$ between them is the path with minimum cost joining them and remaining in $\mathcal{D}$. As for 2-D, it is well known that geodesic paths are simple. The segments in a geodesic path can be of the following types: 1) *cell-crossing* segments which cross a tetrahedra joining two points on its boundary that lie on a different faces, 2) *face-using* segments which lie along a face, and 3) *edge-using* segments which lie along an edge. Consecutive segments of a geodesic path obey Snell's law of refraction in which the path bends at the faces and edges of $\mathcal{D}$. Steiner points in $\mathcal{D}$ will be defined with the aid of the following:

DEFINITION 3.1. *For a point $x \in \mathcal{D}$, define $d(x)$ to be the Euclidean distance from x to the boundary of the union of the tetrahedra incident to x.*

For example, if x lies on a face shared by two tetrahedra, then $\delta(x)$ is the distance from x to the boundary of the union of these two tetrahedra. For a vertex v of $\mathcal{D}$ we define a radius $r(v) = \frac{w_{\min}(v)}{cw_{\max}(v)}d(v)$, where $w_{\min}(v)$ ($w_{\max}(v)$) is the minimum (maximum) weight around v and $c \geq 1$ is a constant. Similarly, for a point x on an edge e (not a vertex) of $\mathcal{D}$ we define a radius $r(x) = \frac{w_{\min}(e)}{cw_{\max}(e)}d(x)$, where $w_{\min}(e)$ ($w_{\max}(e)$) is the minimum (maximum) weights around e. Also we define $r(e)$ to be the maximum of $r(x)$, for $x \in e$, i.e. $r(e) = \max_{x \in e} r(x)$.

DEFINITION 3.2. *A vertex-vicinity $\mathcal{V}(v)$ is the ball, $B(v, \varepsilon r(v))$, centered at v of radius $\varepsilon r(v)$. For an edge $e = (v', v'')$ we define an edge-vicinity $\mathcal{V}(e) = \mathcal{V}(v') \cup \mathcal{V}(v'') \cup \mathcal{S}(e)$, where the spindle $\mathcal{S}(e)$ of e is defined by $\mathcal{S}(e) = \cup_{x \in e} B(x, \varepsilon r(x))$.*

3.3 Placement of Steiner Points
The placement of Steiner points is more complex compared to the placement in the SPPS problem since, a shortest path may bend in the interior of a face. We therefore need to take care of edge vicinities in addition to the vertex-vicinities. An infinitesimal segment may penetrate through these vicinities and any placement of a finite set of Steiner points and edges cannot provide good approximations to these segments. This is addressed by providing an amortization argument for the cost of the part of the path inside these vicinities with respect to the total cost of the path. Let f be a face with vertices A, B and C and let M be the point in f with maximum distance $d(M)$, i.e., $d(M) = \max_{x \in f} d(x)$. The point M is defined to be a Steiner point. We describe the placement of Steiner points inside the triangle ABM. Let MH be the height of the triangle ABM. Define an infinite sequence of points $P_0, P_1, \ldots$ on MH as follows: $P_0 = M$ and for $i = 1, \ldots, |P_{i-1}P_i| = \varepsilon d(P_i)$. Consider the corresponding sequence of segments A_iB_i that are parallel to AB and contain P_i for $i = 1, \ldots$ On each of the segments A_iB_i define a set of $k_i + 1$ equi-distantly placed points $P_{i,j}$, $j = 0, \ldots, k_i$ with

$$k_i = \lceil |A_i, B_i|/(\varepsilon d(P_{i+1})) \rceil. \qquad (24)$$

The formal description of these points is given by $P_{i,0} = A_i$, $P_{i,k_i} = B_i$ and for $j = 1, \ldots, k_i$, $|P_{i,j-1}P_{i,j}| = |A_iB_i|/k_i$.

DEFINITION 3.3. *The set of Steiner points in the triangle ABM consists of all points $P_{i,j}$ that are outside the edge-vicinity $\mathcal{V}(AB)$.*

The total number of Steiner points placed in $\mathcal{D}$ is given by the following lemma whose proof can be found in the full version [3].

LEMMA 3.1. *The number of Steiner points placed in a face f of $\mathcal{D}$ does not exceed $C_f(\mathcal{D})\frac{1}{\varepsilon^3}\ln\frac{1}{\varepsilon}$, where the constant $C_f(\mathcal{D})$ depends on geometric features of $\mathcal{D}$ around f. The total number of Steiner points is bounded by $C(\mathcal{D})\frac{n}{\varepsilon^3}\ln\frac{1}{\varepsilon}$, where n is the number of faces of $\mathcal{D}$ and the constant $C(\mathcal{D})$ is the arithmetic average of the constants $C_f(\mathcal{P})$ for the faces of $\mathcal{D}$.*

3.4 Forming the approximation graph

Consider a segment ab of a geodesic path lying inside a trapezoid t_i, such that a and b belong to two different faces f_1 and f_2 of t_i, respectively. Assume that ab does not intersect any vertex or edge vicinities. It can be shown that there exist Steiner points $a' \in f_1$ and $b' \in f_2$ so that weighted length of the path from a to b through a' and b' is at most $1 + \varepsilon$ times the cost of the direct path from a to b. This enables us to provide ε approximations to the part of the geodesic path that does not lie inside vertex or edge vicinities. The part of the geodesic path that lies inside vertex-vicinities can be handled in a similar manner as in the SPPS problem, i.e., amortizing its cost with respect to the cost of the path that is outside the vertex vicinities. Now concentrate on the part of the path (say $\Pi(e)$) that lies inside an edge vicinity $\mathcal{V}(e)$. There are two cases depending on the cost of $\Pi(e)$ with respect to the parameter $\varepsilon r(x)$ (Definition 3.2). If the two costs are comparable then use the same amortization argument as in the case of vertex-vicinities. If the cost of $\Pi(e)$ is more than $\varepsilon r(x)$ then there exist a pair of Steiner points a' and b' on the boundary of $\mathcal{V}(e)$, such that the cost of the path between these Steiner points is at most $1 + \varepsilon$ times the cost of $\Pi(e)$.

The above discussion suggests that if we interconnect all Steiner points within a trapezoid, for every trapezoid in $\mathcal{D}$, and interconnect all Steiner points on the boundary of edge vicinities, then we can obtain a graph G, where vertices are Steiner points, and the edges are the interconnections. For any linear path $\tilde{\Pi}$ joining a pair of vertices of $\mathcal{D}$, it follows that there exists a discrete path Π between the corresponding vertices in G whose cost is at most $1 + \varepsilon$ times the cost $\tilde{\Pi}$. Although the approximation quality of G is the desired one, the number of edges, and hence the computational complexity of finding a shortest path in G, is very high. Next we show that using properties of Snell's law, the search can be restricted to a selected set of edges when a vertex is "explored" (after extract-min) in an execution of Dijkstra's algorithm.

3.5 Pruned Dijkstra's algorithm

Let a be a bending point on a geodesic path Π lying in the interior of a face f. The in-angle φ is defined to be the acute clockwise angle between the normal to f at a and the segment s' preceding a on Π. Similarly, the out-angle ψ is the acute clockwise angle between the normal and the segment s'' succeeding a on Π. These angles are related by $w'' \sin \psi = w' \sin \varphi$, where the cost of traveling along s' and

s'' is w' and w'', respectively. Moreover the segments s', s'', and the normal to f at a (denoted by $\mathcal{N}_{fb}$) lie in the same plane. This can be expressed by taking a reference direction in f, and the angles $\eta_{in} = \eta_{out}$, where η_{in} (η_{out}) is the angle between the projection of s' (s'') on f and the reference direction. Next we estimate the difference between out-angles of two given geodesic paths in terms of the difference between their in-angles.

LEMMA 3.2. *Let $(\varphi, \psi, \eta_{in}, \eta_{out})$ and $(\varphi_1, \psi_1, \eta_{1in}, \eta_{1out})$, be two pairs of angles satisfying the above equations. Then*

$$(a) \quad |\psi_1 - \psi| \leq \frac{\pi w'}{2w''}\kappa(\varphi, \varphi_1)|\varphi_1 - \varphi| \qquad (25)$$

$$(b) \quad |\eta_{1in} - \eta_{in}| = |\eta_{1out} - \eta_{out}| \qquad (26)$$

where $\kappa(\varphi, \varphi_1) = \cos\frac{\varphi_1+\varphi}{2} / \cos\frac{\psi_1+\psi}{2}$.

The following lemma estimates the dispersion of a discrete path with respect to the geodesic path (analogous to Lemma 2.6).

LEMMA 3.3. *Let $\varepsilon < 1$ and a, b and c be three consecutive bending points on a discrete path Π, where the Steiner point b lying in the face f does not neighbour a vertex or edge vicinity of $\mathcal{D}$. Let c_0 be a point on the geodesic direction determined by the segment (a, b) (on the ray starting at b and with an out-angle ψ, η_{out} satisfying Snell's law). Then, (a) the angle between the two planes determined by $\mathcal{N}_{fb}$ and bc_0 and bc is at most ε. (b) the angle between the projection of bc, onto the plane defined by the bc_0 and $\mathcal{N}_{fb}$, and bc_0 is at most $c'\sqrt{\varepsilon}$, where c' is a constant.*

The above lemma establishes that during an execution of Dijkstra's algorithm, it suffices to examine only those edges which are given by the specified angle constraint. Following along the lines of Section 2.4, it can be shown that for all Steiner points which are not neighboring vertex or edge vicinities, we have to explore only $O(1/\sqrt{\varepsilon})$ edges. We summarize the result in the following theorem:

THEOREM 3.1. *Let $\mathcal{D}$ be a weighted triangulation of 3-dimensional space and $0 < \varepsilon < 1$ consisting of n vertices. The single source shortest paths on $\mathcal{D}$ can be approximated to within an accuracy of $(1 + c\varepsilon)$ in $O(\frac{n}{\varepsilon^3}\log\frac{1}{\varepsilon}(\frac{1}{\sqrt{\varepsilon}} + \log n))$ time, where $c \geq 1$ is a constant.*

In case all the weights are the same, the complexity reduces to $O(\frac{n}{\varepsilon^3}\log\frac{1}{\varepsilon}\log n)$ time. This can used to solve the shortest path problem amidst obstacles in 3-dimensional Euclidean space (ESP-3D), provided that we are given a convex partition (or tetrahedralization) of the free space. Convex partition can be achieved by employing any one of the algorithms of [14; 4; 11].

4. REFERENCES

[1] K.R. Varadarajan and P.K. Agarwal, "Approximating Shortest Paths on a Polyhedron", *38FOCS*, 1997.

[2] L. Aleksandrov, M.Lanthier, A. Maheshwari, J.-R. Sack, "An ε-Approximation Algorithm for Weighted Shortest Paths", Proc. SWAT, Stockholm, LNCS 1432, July 1998, pp. 11–22.

[3] L. Aleksandrov, A. Maheshwari and J.-R. Sack, "Approximation Algorithms for Geometric Shortest Path Problems", technical report, School of Compter Science, TR-99-10, October 1999.

[4] C. Bajaj and T. Dey, "Convex decompositions of polyhedra and robustness", SIAM Jl. Comput. 21, 1992, pp. 339-364.

[5] J. Canny and J. H. Reif, "New Lower Bound Techniques for Robot Motion Planning Problems", *28th IEEE Symp. on Foundations of Computer Science*, 1987, pp. 49-60.

[6] J. Chen and Y. Han, "Shortest Paths on a Polyhedron", *6th SoACM-CG*, 1990, pp. 360-369.

[7] K.L. Clarkson, "Approximation algorithms for shortest path motion planning", *19 ACM Symposium on Theory of Computing*, 1987, pp. 56-65.

[8] J. Choi, J. Sellen and C.K. Yap, "Approximate Euclidean Shortest Path in 3-Space", *10 SoACM-CG*, 1994, pp. 41-48.

[9] J. Choi, J. Sellen and C.K. Yap, "Approximate Euclidean Shortest Path in 3-Space", *Intl. Jl. of Computational Geometry and Applications*,Vol. 7, no. 4, 1997, pp. 271–295.

[10] G. Das and G. Narasimhan, "Short Cuts in Higher Dimensional Space", *Proceedings of the 7th Annual Canadian Conference on Computational Geometry*, Québec City, Québec, 1995, pp. 103-108.

[11] T. Dey, "personal communication", March 1999.

[12] E.W. Dijkstra, "A Note on Two Problems in Connection with Graphs", *Numerical Mathematics 1*, 1959, pp.269-271.

[13] M.L. Fredman and R.E. Tarjan, "Fibonacci Heaps and Their Uses in Improved Network Optimization Algorithms", *J. ACM*, **34**(3), 1987, pp.596-615.

[14] J. Hershberger and J. Snoeyink, "Erased arrangements of lines and convex decompositions of polyhedra", CGTA 9, 1998, pp. 129–143.

[15] S. Kapoor, "Efficiently computing geodesic shortest paths", 1999 ACM Symposium on Theory of Computing.

[16] C. Kenyon and R. Kenyon, "How To Take Short Cuts", *Discrete and Computational Geometry*, Vol. 8, 1992, pp. 251-264.

[17] M. Lanthier, A. Maheshwari and J.-R. Sack, "Approximating Weighted Shortest Paths on Polyhedral Surfaces", *13 SoACM-CG*, 1997, pp. 274-283. (to appear in Algorithmica)

[18] M. Lanthier, "Shortest Path Problems on Polyhedral Surfaces", *Ph.D. Thesis* , School of Computer Science, Carleton University, Ottawa, Canada, 2000.

[19] M. Lanthier, A. Maheshwari, and J.-R. Sack, "Shortest Anisotropic Paths on Terrains", in Proc. ICALP'99, Prague, LNCS 1644, pp. 524-533, 1999.

[20] C. Mata and J. Mitchell, "A New Algorithm for Computing Shortest Paths in Weighted Planar Subdivisions", *13 SoACM-CG*, 1997, pp. 264-273.

[21] J.S.B. Mitchell and C.H. Papadimitriou, "The Weighted Region Problem: Finding Shortest Paths Through a Weighted Planar Subdivision", *Journal of the ACM*, **38**, January 1991, pp. 18-73.

[22] C.H. Papadimitriou, "An Algorithm for Shortest Path Motion in Three Dimensions", *IPL*, **20**, 1985, pp. 259-263.

[23] M. Sharir, "On Shortest Paths Amidst Convex Polyhedra", *SIAM Jl. Comp.*, **16**, 1987, pp. 561-572.